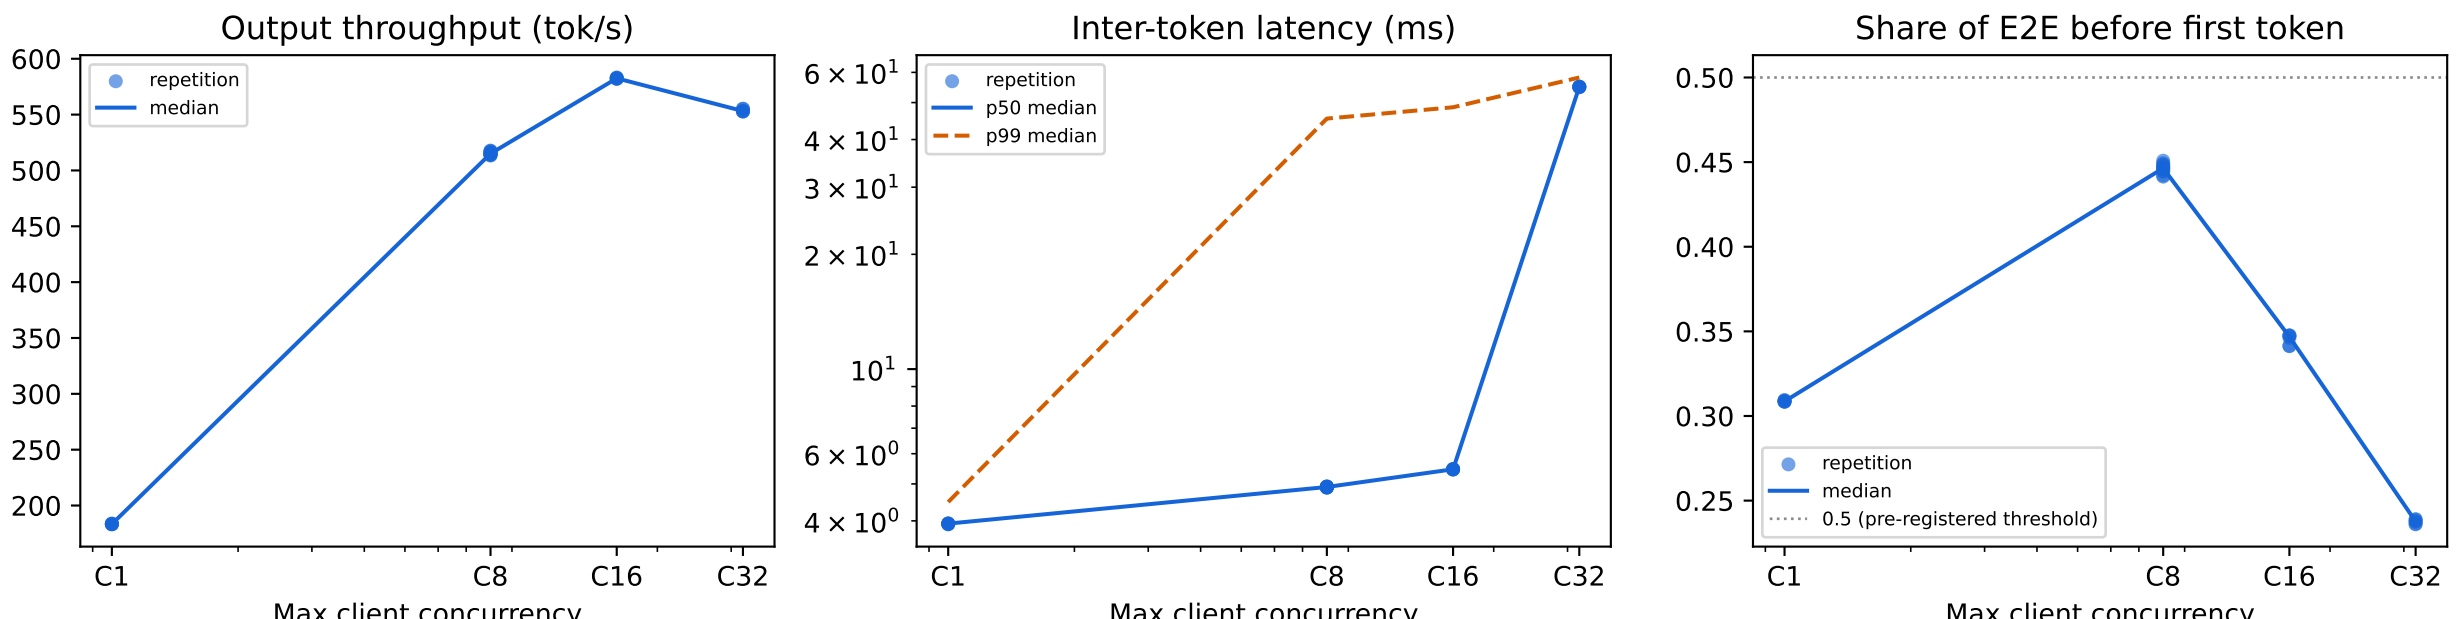
prefill_heavy (2048 in / 32 out): throughput saturates by C16, and the extra latency moves into per-step time
Output throughput (tok/s)
600
550
500
450
400
350
300
250
200
repetition
median
C1
C8
C16
C32
Max client concurrency
Inter-token latency (ms)
6×10^1
4×10^1
3×10^1
2×10^1
10^1
6×10^0
4×10^0
repetition
p50 median
p99 median
C1
C8
C16
C32
Max client concurrency
Share of E2E before first token
0.50
0.45
0.40
0.35
0.30
0.25
repetition
median
0.5 (pre-registered threshold)
C1
C8
C16
C32
Max client concurrency