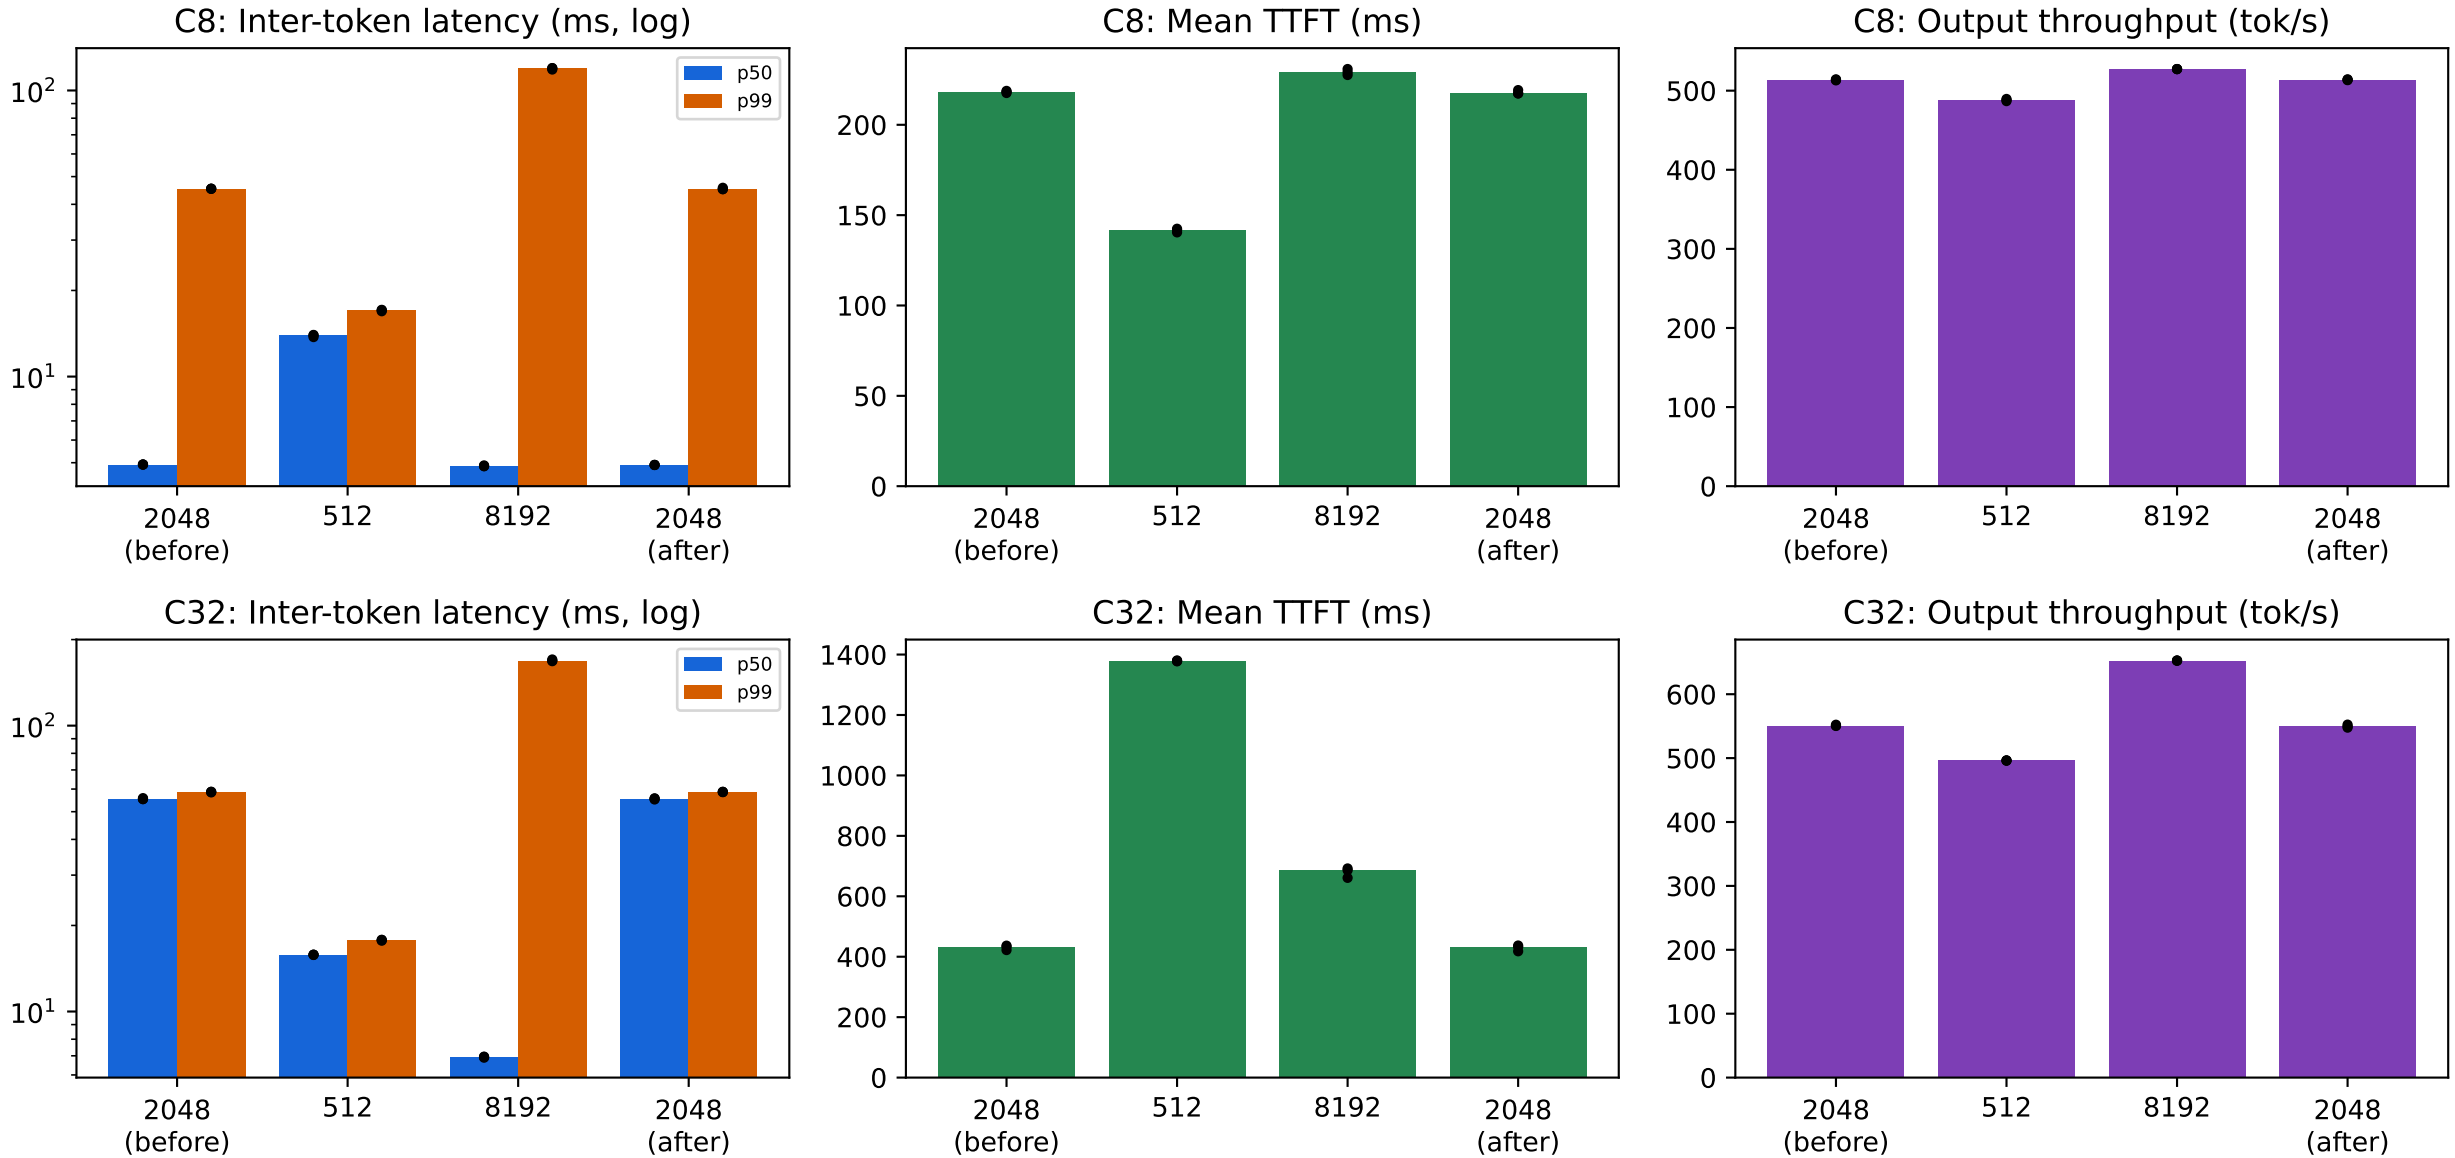
Phase D token budget (max_num_batched_tokens): small budgets give frequent small stalls, large budgets rare large ones
C8: Inter-token latency (ms, log)
p50
p99
10^2
10^1
2048 (before)
512
8192
2048 (after)
C8: Mean TTFT (ms)
0
50
100
150
200
C8: Output throughput (tok/s)
0
100
200
300
400
500
C32: Inter-token latency (ms, log)
C32: Mean TTFT (ms)
0
200
400
600
800
1000
1200
1400
C32: Output throughput (tok/s)
0
100
200
300
400
500
600