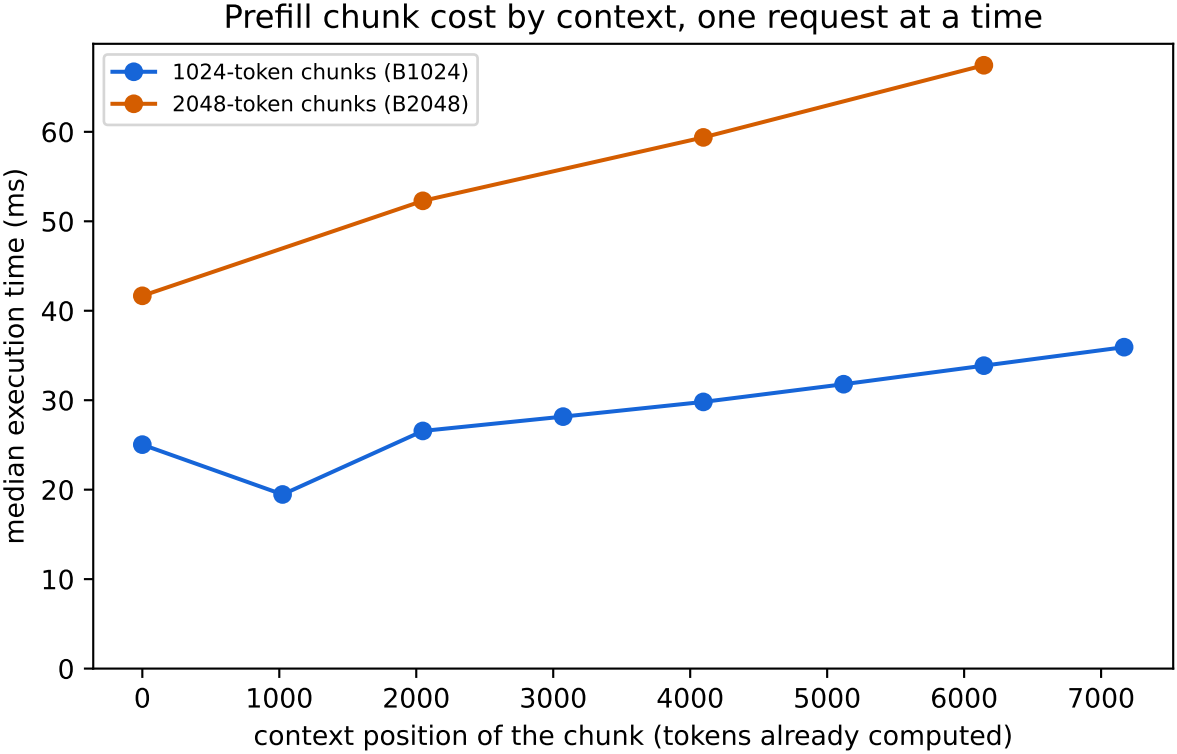
Prefill chunk cost by context, one request at a time
1024-token chunks (B1024)
2048-token chunks (B2048)
median execution time (ms)
0
10
20
30
40
50
60
0
1000
2000
3000
4000
5000
6000
7000
context position of the chunk (tokens already computed)